

## Torch Library loading

```
In [28]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader, random_split
import matplotlib.pyplot as plt
import numpy as np
import math
from torch.autograd import Variable
from torchvision.utils import make_grid
import os
from torchsummary import summary
```

## Device configuration

```
In [2]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
device
Out[2]: device(type='cuda')
```

## Hyper-parameters

```
In [3]: num_epochs = 15
batch_size = 32
learning_rate = 0.01
```

## Data loading and downloading

```
In [4]: # dataset has PILImage images of range [0, 1].
# We transform them to Tensors of normalized range [-1, 1]
transform = transforms.Compose([
    transforms.ToTensor()
])

# CIFAR10: 60000 32x32 color images in 10 classes, with 6000 images per class
train_dataset = torchvision.datasets.MNIST(root='./data', train=True,
download=True, transform=transform)
test_dataset = torchvision.datasets.MNIST(root='./data', train=False,
download=True, transform=transform)
TRAIN_SPLIT=0.9
VAL_SPLIT=0.1
numTrainSamples = int(len(train_dataset) * TRAIN_SPLIT)
numValSamples = int(len(train_dataset) * VAL_SPLIT)
(trainData, valData) = random_split(train_dataset,
(numTrainSamples, numValSamples),
generator=torch.Generator().manual_seed(42))
train_loader = torch.utils.data.DataLoader(trainData, batch_size=batch_size,
shuffle=True)
validation_loader=torch.utils.data.DataLoader(valData, batch_size=batch_size,
shuffle=True)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size,shuffle=False)
```

## imshow()

```
In [5]: classes = ('0', '1', '2', '3',
'4', '5', '6', '7', '8', '9')
def imshow(img):
    img = img / 2 + 0.5 # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()
```

## 1 MNIST classification using CNN

```
In [6]: class CNN(nn.Module):
def __init__(self):
    super(CNN, self).__init__()
    self.conv1 = nn.Conv2d(1,32,kernel_size = 3, stride = 1, padding = 1)
    self.conv2 = nn.Conv2d(32,32,kernel_size = 3, stride = 1, padding = 1)
    self.fc1 = nn.Linear(7*7*32, 500)
    self.fc2 = nn.Linear(500, 10)
    self.activ = nn.ReLU()

    def pool(self, x, kernel_size = 2, stride = 2):
        out = F.max_pool2d(x, kernel_size, stride)
        return out

    def forward(self, x, softmax = True):
        out = self.activ(self.conv1(x))
        out = self.pool(out)
        out = self.activ(self.conv2(out))
        out = self.pool(out)
        out = out.reshape(out.size(0),-1)
        out = self.activ(self.fc1(out))
        out = self.fc2(out)
        if softmax:
            return F.softmax(out, dim = 1)
        else:
            return out

In [7]: CNN().forward

Out[7]: <bound method CNN.forward of CNN(
  (conv1): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (fc1): Linear(in_features=1568, out_features=500, bias=True)
  (fc2): Linear(in_features=500, out_features=10, bias=True)
  (activ): ReLU())
>
```

## Training Error ,Validation error plots for every epoch and average prediction accuracy

```
In [8]: # get some random training images
dataiter = iter(train_loader)
images, labels = next(dataiter)
# show images
imshow(torchvision.utils.make_grid(images))
model = CNN().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
n_total_steps = len(train_loader)
training_loss=[]
validation_loss=[]
n_correct = 0
n_class_correct = 0
for i in range(10):
    n_class_samples = 0
    for i in range(10):
        train_loss_epoch=[]
        validation_loss_epoch=[]
        num_iter=int(math.ceil(len(trainData)/batch_size))
        for epoch in range(num_epochs):
            training_loss=[]
            validation_loss=[]
            for i in range(num_iter):
                images,labels=next(iter(train_loader))
                images_val,labels_val=next(iter(validation_loader))

                images = images.to(device)
                labels = labels.to(device)
                images_val=images_val.to(device)
                labels_val=labels_val.to(device)

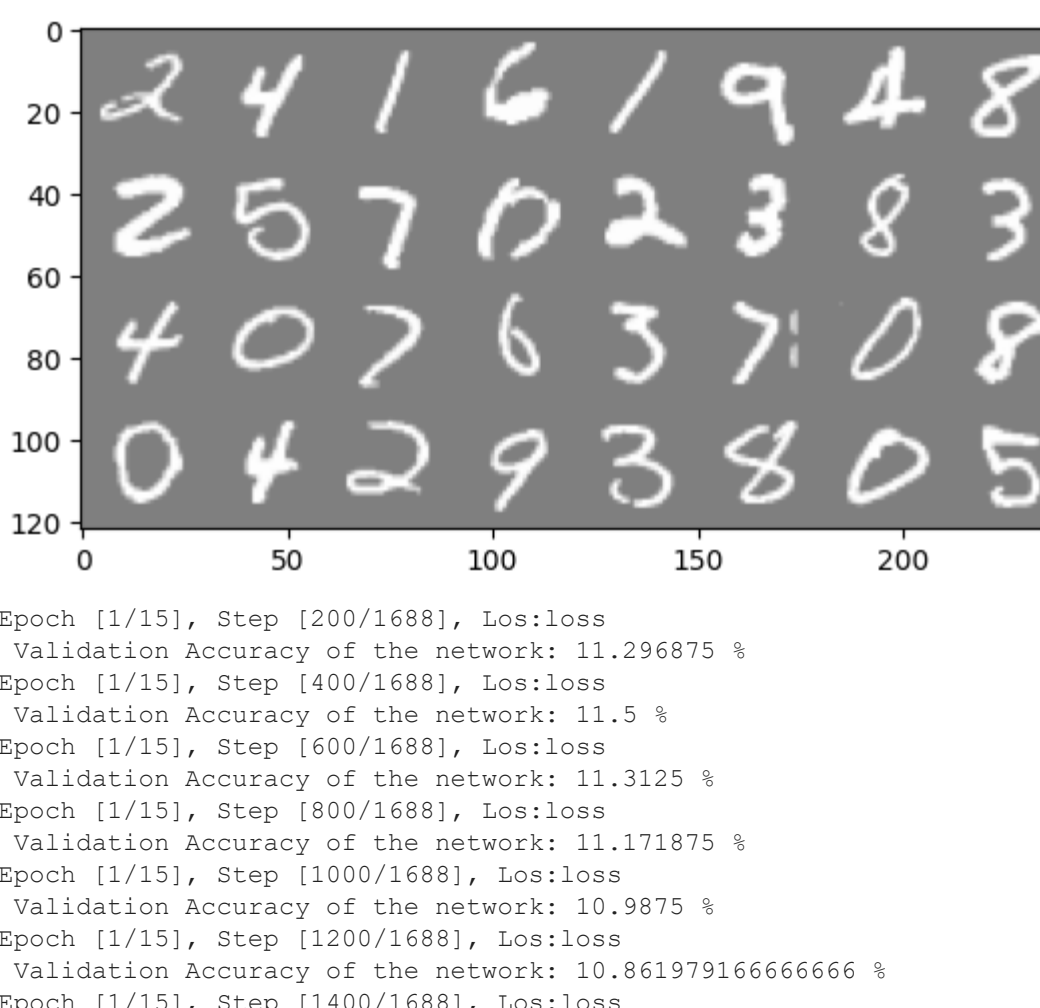
                # Forward pass
                outputs = model(images)
                outputs_val=model(images_val)
                loss = criterion(outputs, labels)

                loss_val=criterion(outputs_val,labels)

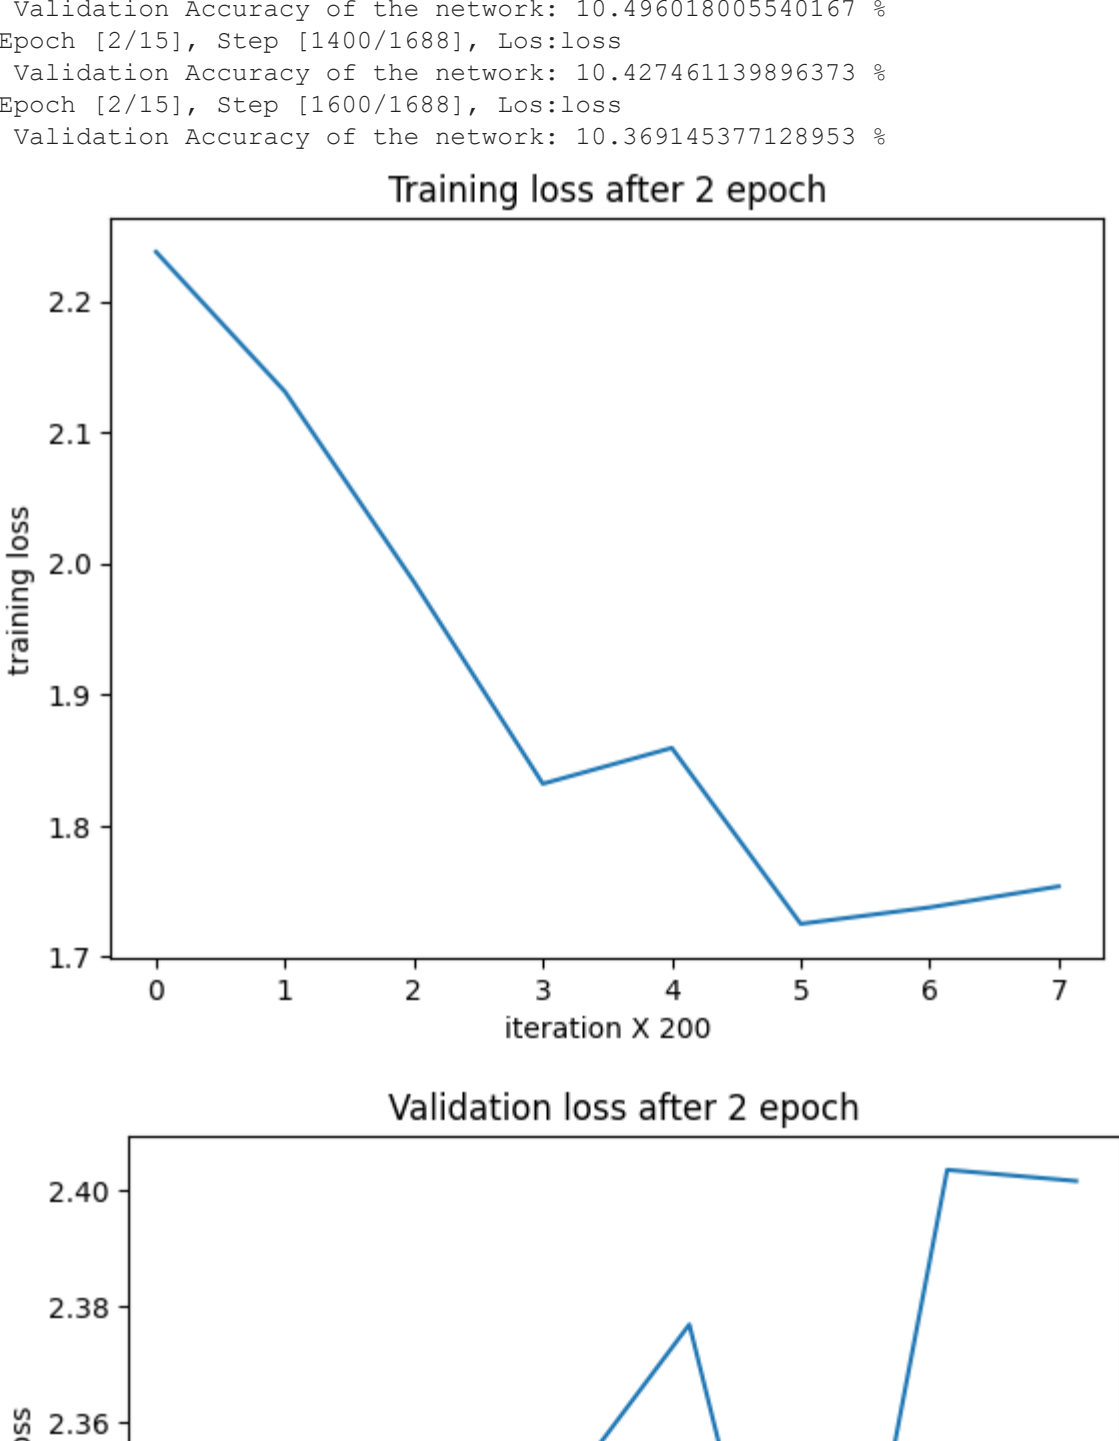

                # Backward and optimize
                optimizer.zero_grad()
                loss.backward()
                optimizer.step()
                train_loss=loss.cpu().detach().numpy()
                val_loss=loss_val.cpu().detach().numpy()

                _, predicted = torch.max(outputs_val, 1)
                n_samples += labels.size(0)
                n_correct += (predicted == labels).sum().item()
                batch_size_list=list(labels.size())
                batch_size=batch_size_list[0]
                for j in range(batch_size):
                    label = labels[j]
                    pred = predicted[j]
                    if (label == pred):
                        n_class_correct[label] += 1
                        n_class_samples[label] += 1
                acc = 100.0 * n_correct / n_samples
                if (i+1) % 200 == 0:
                    print ('Epoch [{epoch+1}/{num_epochs}], Step [{i+1}/{n_total_steps}], Loss:loss')
                    print (' Validation Accuracy of the network: (acc) %')
                    training_loss.append(train_loss)
                    validation_loss.append(val_loss)
            plt.plot(range(len(training_loss)),training_loss)
            plt.xlabel("Iteration X 200")
            plt.ylabel("training loss")
            plt.title("Training loss after (epoch+1) epoch")
            plt.show()
            plt.plot(range(len(validation_loss)),validation_loss)
            plt.xlabel("Iteration X 200")
            plt.ylabel("Validation loss")
            plt.title("Validation loss after (epoch+1) epoch")
            plt.show()
            train_loss_epoch.append(train_loss)
            validation_loss_epoch.append(val_loss)
        fig,ax=plt.subplots()
        ax2=ax1.twinx()
        line1=ax1.plot(range(len(train_loss_epoch)),train_loss_epoch,color="blue",label="trainin loss")
        ax1.set_ylabel("training loss")
        line2=ax2.plot(range(len(train_loss_epoch)),validation_loss_epoch,color="orange",label="validation loss")
        ax2.set_ylabel("validation loss")
        lines=line1+line2
        labels=[l.get_label() for l in lines]
        ax1.legend(lines,labels)
        print('Finished Training')
        PATH = './cnn.pth'
        torch.save(model.state_dict(), PATH)
```


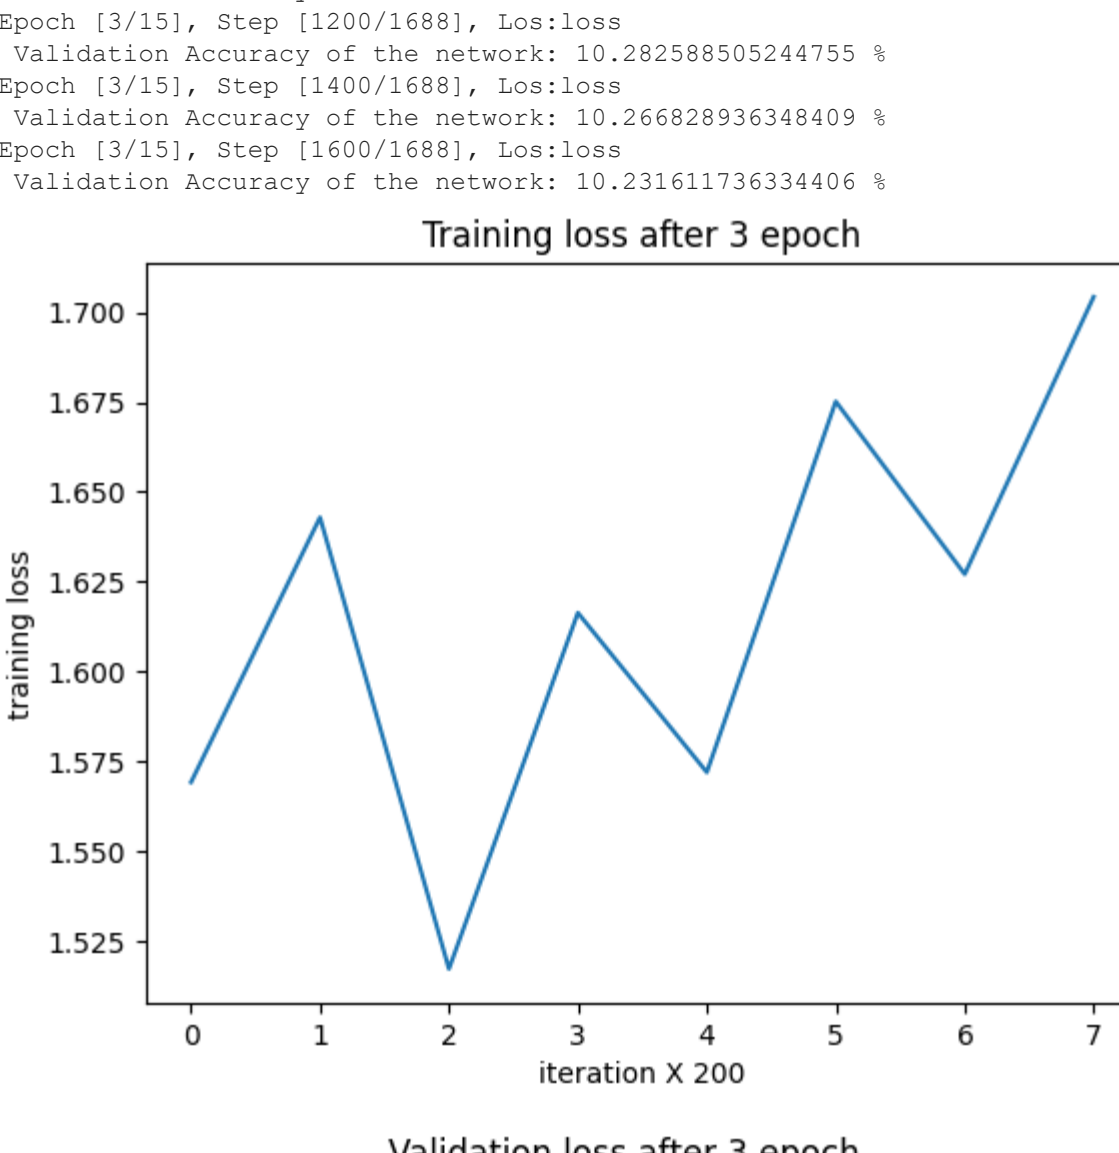




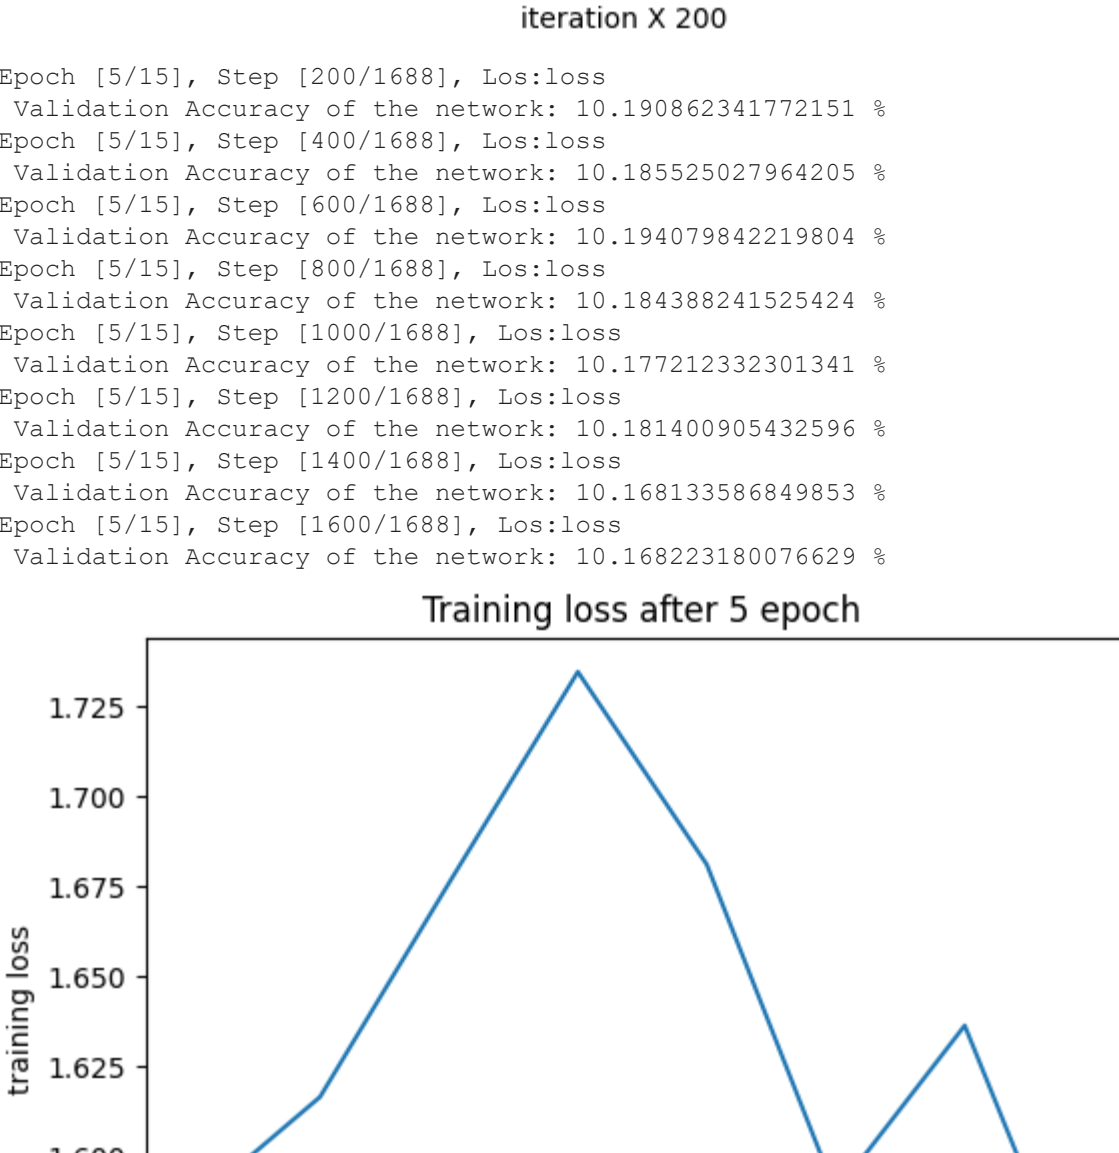
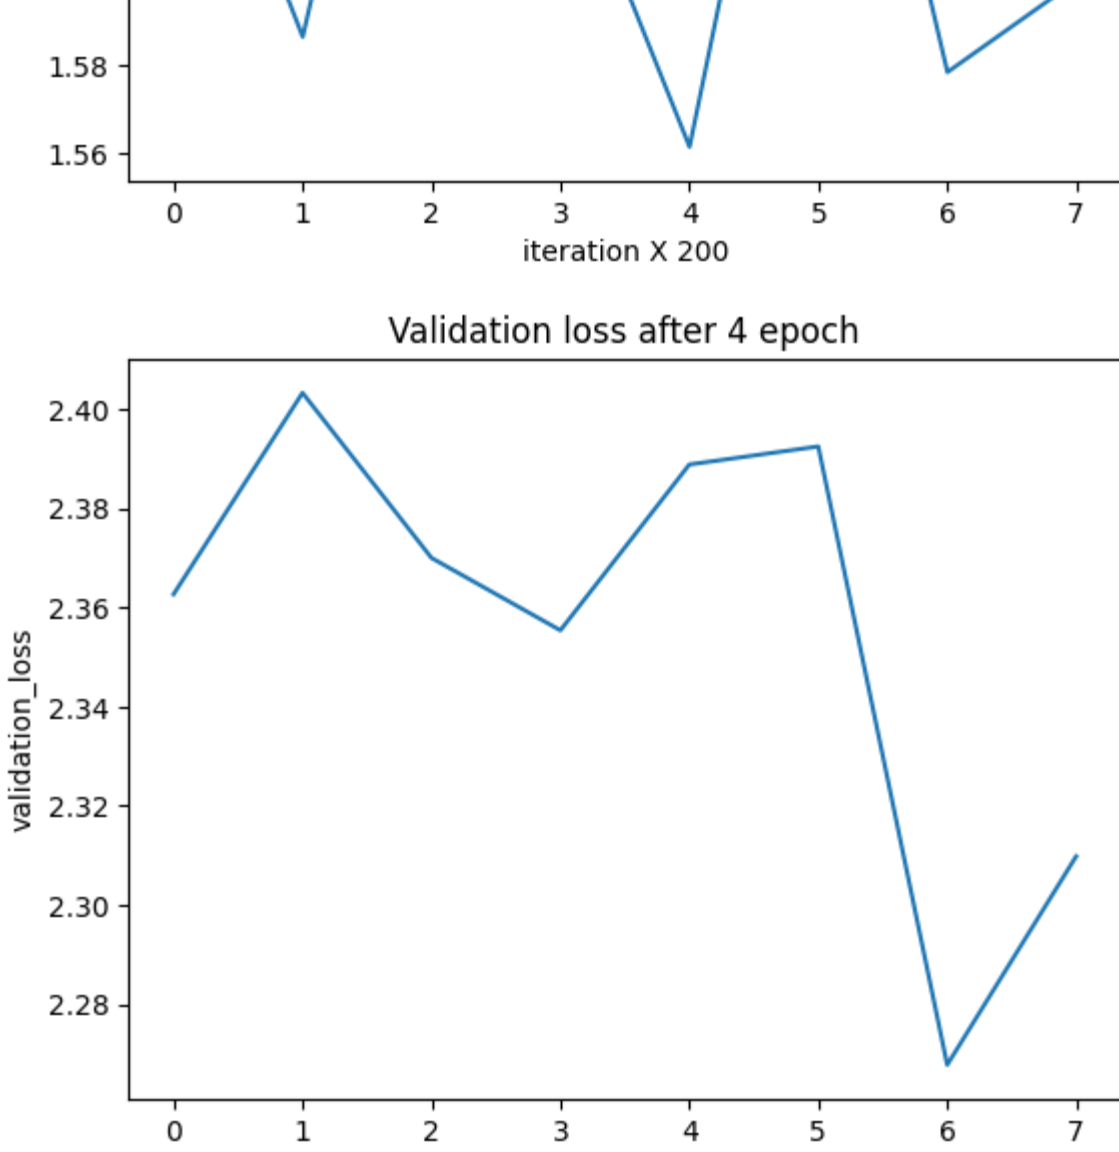
Epoch [1/15], Step (200/1688), Loss:loss  
Validation Accuracy of the network: 11.296875 %  
Epoch [1/15], Step (400/1688), Loss:loss  
Validation Accuracy of the network: 11.5 %  
Epoch [1/15], Step (600/1688), Loss:loss  
Validation Accuracy of the network: 11.3125 %  
Epoch [1/15], Step (800/1688), Loss:loss  
Validation Accuracy of the network: 11.17875 %  
Epoch [1/15], Step (1000/1688), Loss:loss  
Validation Accuracy of the network: 10.9875 %  
Epoch [1/15], Step (1200/1688), Loss:loss  
Validation Accuracy of the network: 10.86197916666666 %  
Epoch [1/15], Step (1400/1688), Loss:loss  
Validation Accuracy of the network: 10.712035371428571 %  
Epoch [1/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.61171875 %



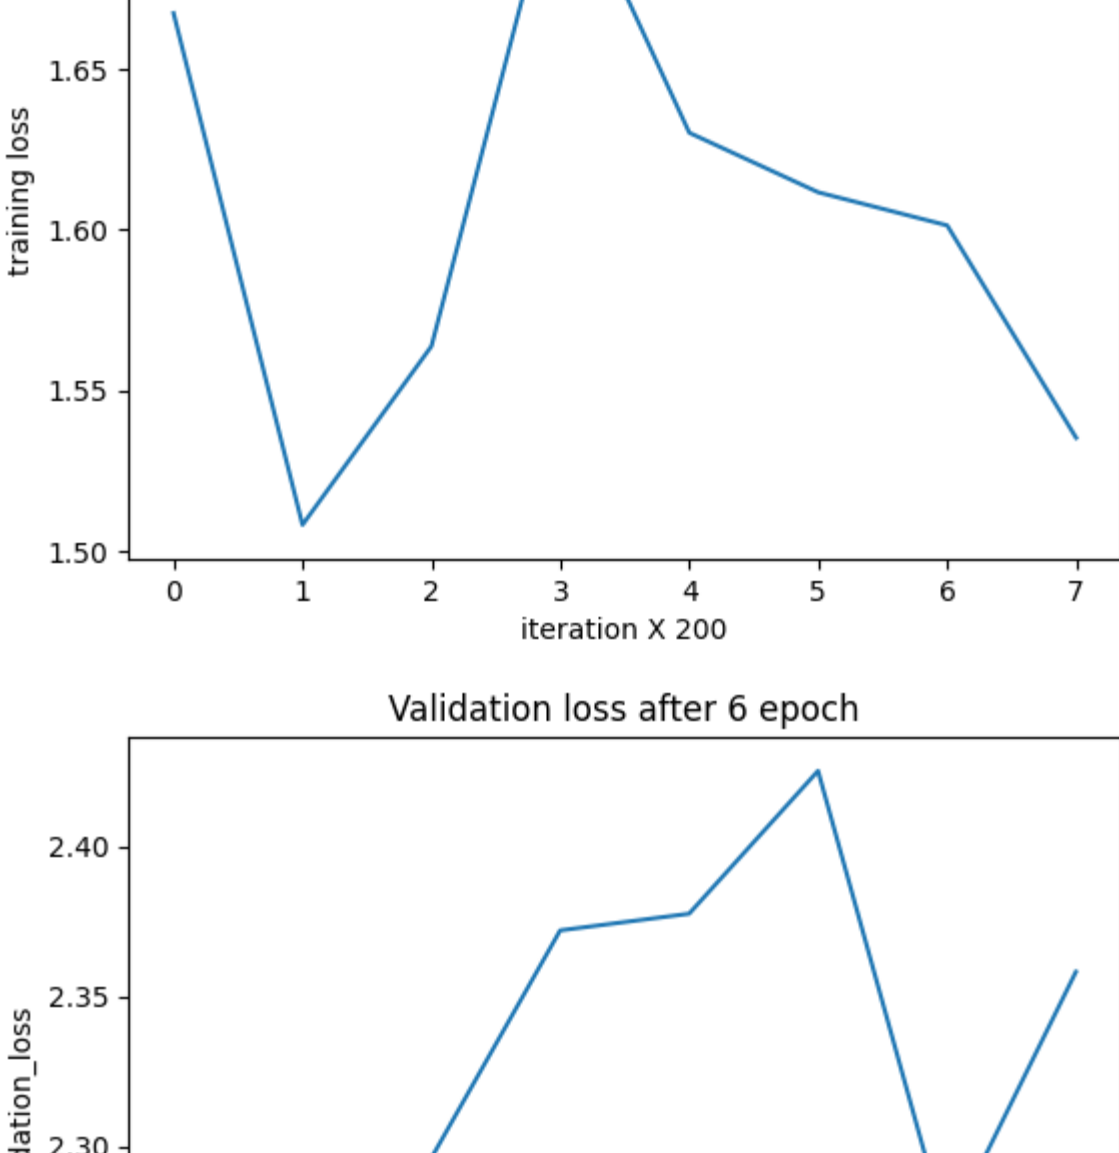
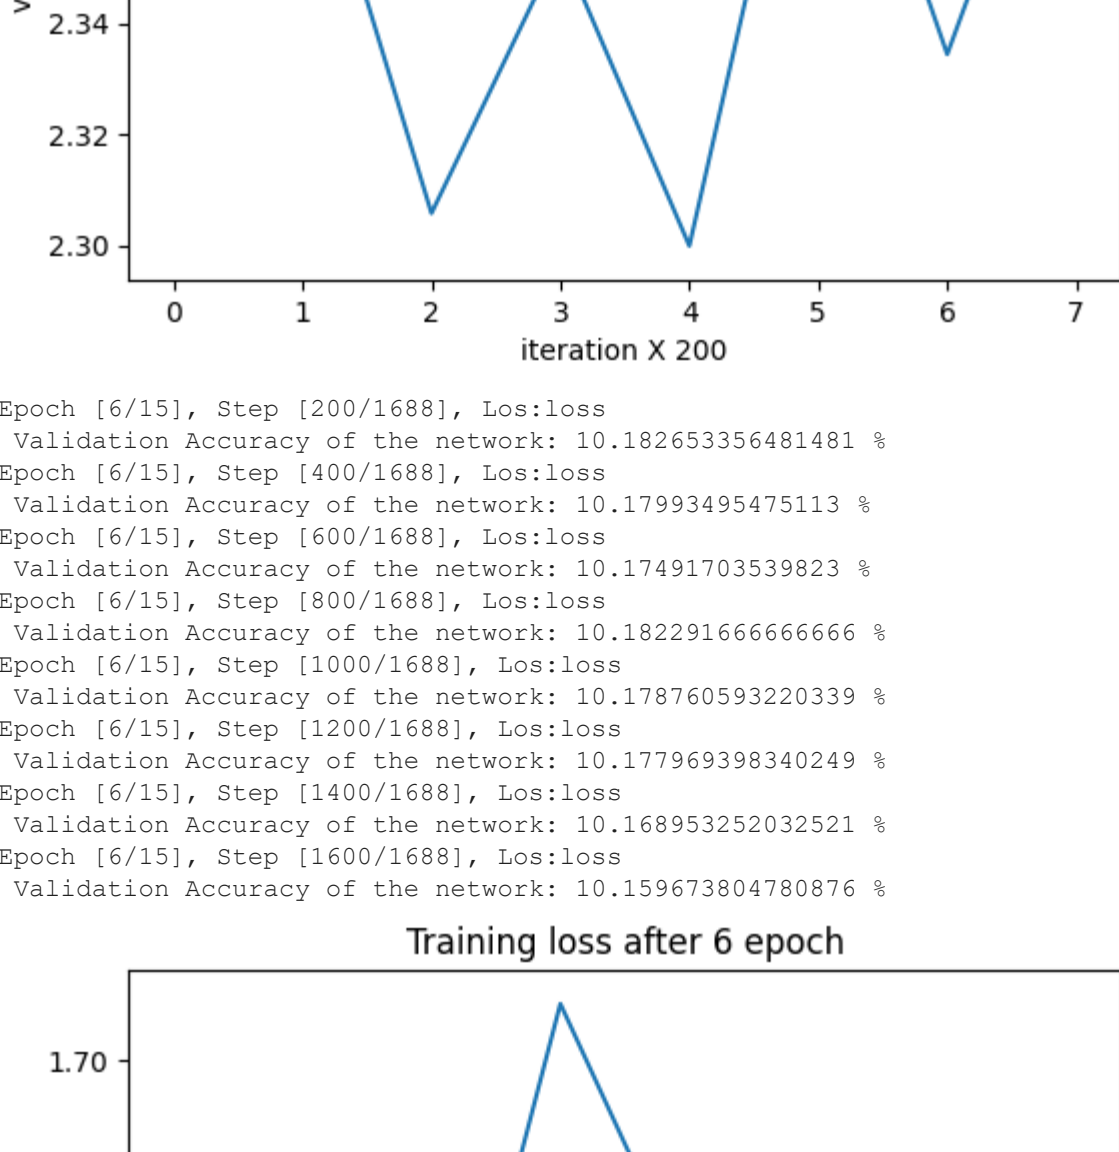
Epoch [2/15], Step (200/1688), Loss:loss  
Validation Accuracy of the network: 10.35025167788135 %  
Epoch [2/15], Step (400/1688), Loss:loss  
Validation Accuracy of the network: 10.49892279693466 %  
Epoch [2/15], Step (600/1688), Loss:loss  
Validation Accuracy of the network: 10.494973776223777 %  
Epoch [2/15], Step (800/1688), Loss:loss  
Validation Accuracy of the network: 10.49030546237942 %  
Epoch [2/15], Step (1000/1688), Loss:loss  
Validation Accuracy of the network: 10.49336577380953 %  
Epoch [2/15], Step (1200/1688), Loss:loss  
Validation Accuracy of the network: 10.49608095846147 %  
Epoch [2/15], Step (1400/1688), Loss:loss  
Validation Accuracy of the network: 10.427461139896373 %  
Epoch [2/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.369145377128953 %




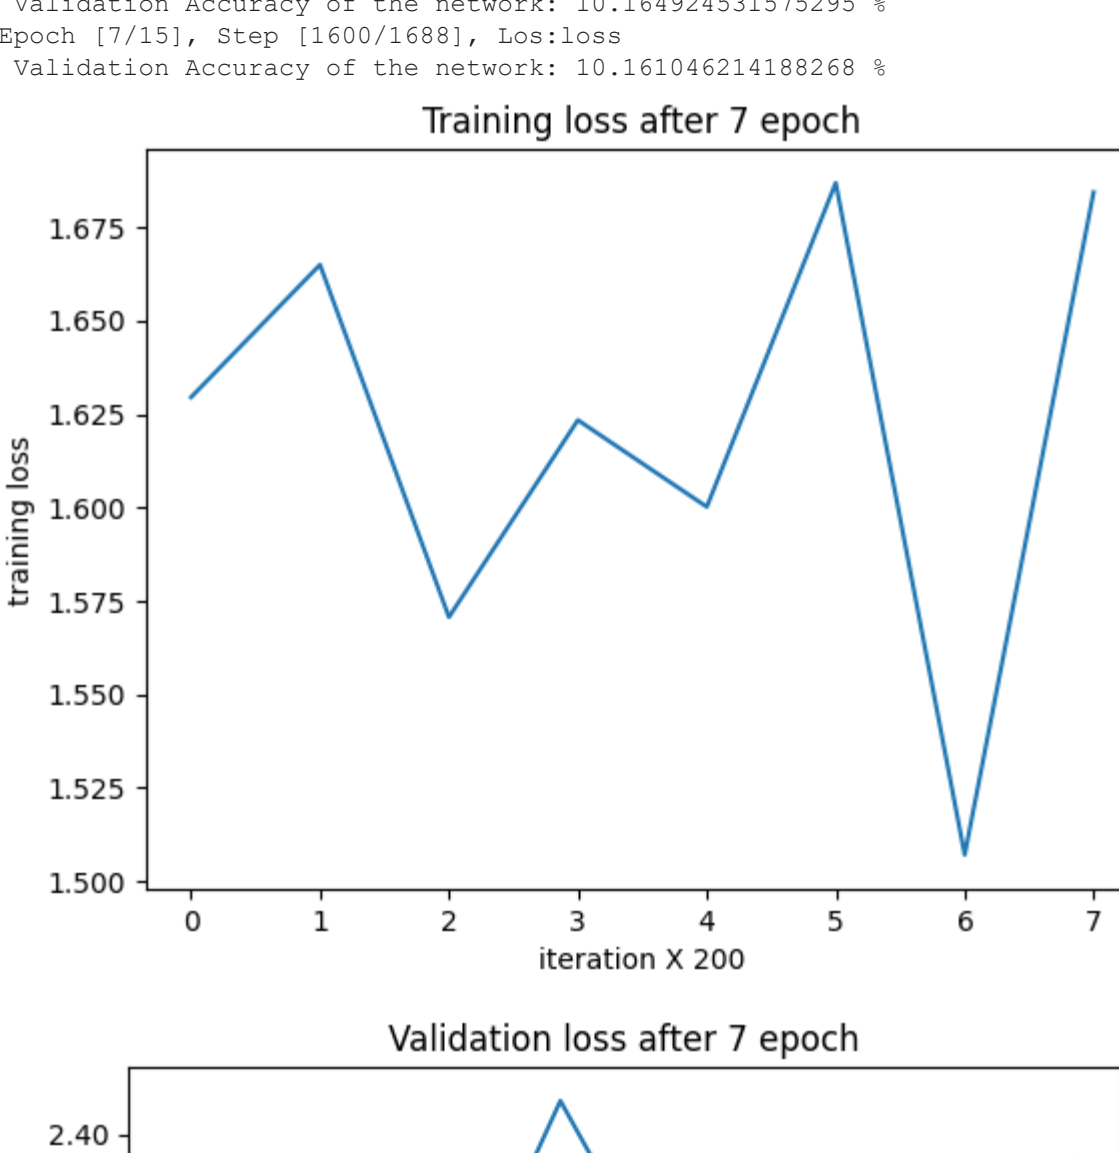
Epoch [3/15], Step (200/1688), Loss:loss  
Validation Accuracy of the network: 10.35025167788135 %  
Epoch [3/15], Step (400/1688), Loss:loss  
Validation Accuracy of the network: 10.3217890679661 %  
Epoch [3/15], Step (600/1688), Loss:loss  
Validation Accuracy of the network: 10.26139614968813 %  
Epoch [3/15], Step (800/1688), Loss:loss  
Validation Accuracy of the network: 10.283465038314176 %  
Epoch [3/15], Step (1000/1688), Loss:loss  
Validation Accuracy of the network: 10.28628007312614 %  
Epoch [3/15], Step (1200/1688), Loss:loss  
Validation Accuracy of the network: 10.282589585244755 %  
Epoch [3/15], Step (1400/1688), Loss:loss  
Validation Accuracy of the network: 10.26682893648409 %  
Epoch [3/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.23161173634406 %



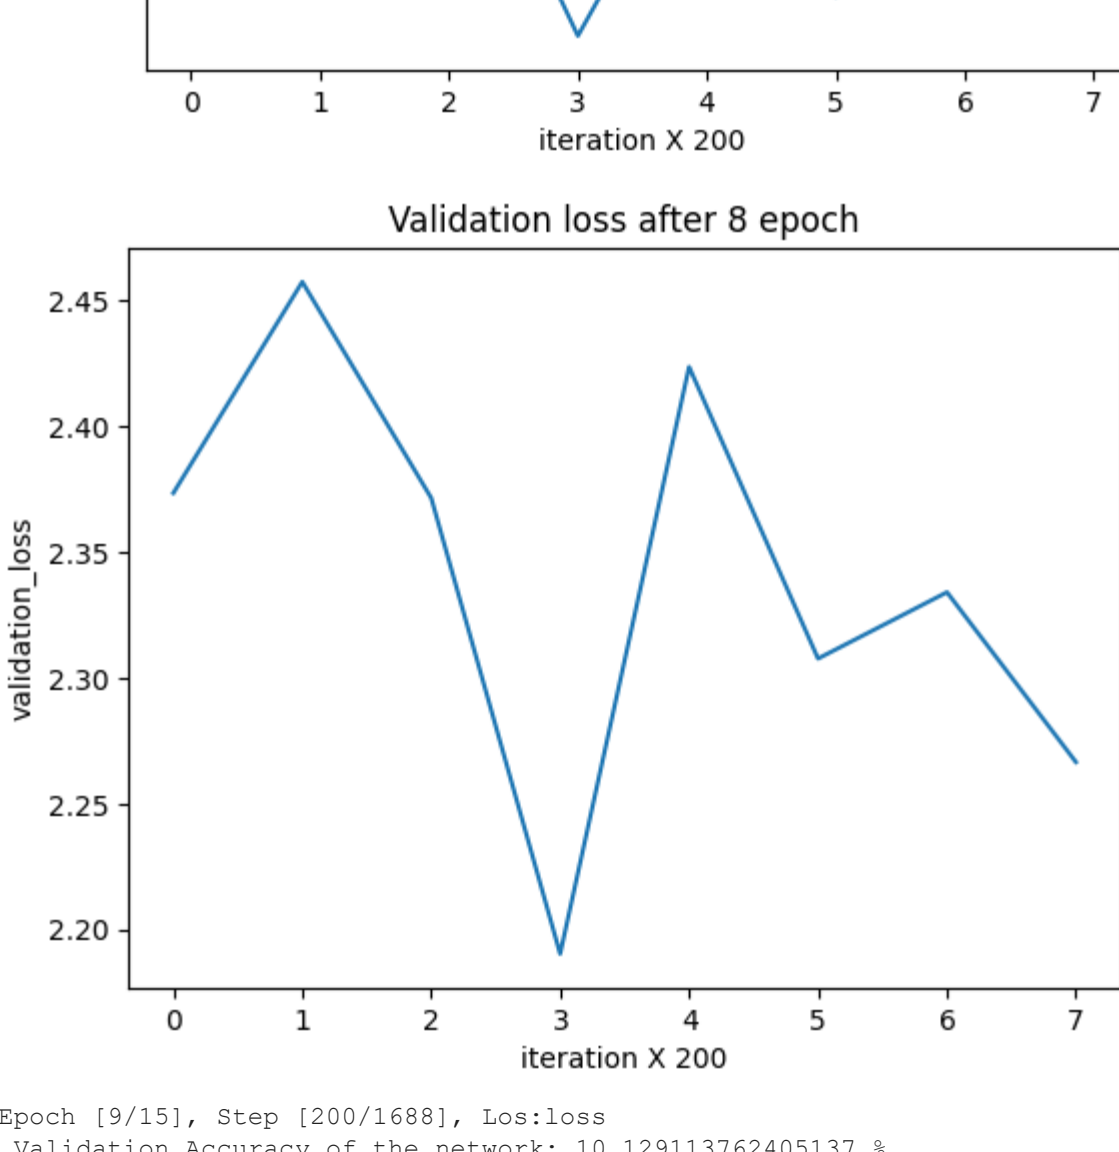
Epoch [4/15], Step (200/1688), Loss:loss  
Validation Accuracy of the network: 10.321644166930092 %  
Epoch [4/15], Step (400/1688), Loss:loss  
Validation Accuracy of the network: 10.1528514140079 %  
Epoch [4/15], Step (600/1688), Loss:loss  
Validation Accuracy of the network: 10.22826405948158 %  
Epoch [4/15], Step (800/1688), Loss:loss  
Validation Accuracy of the network: 10.2373718887262 %  
Epoch [4/15], Step (1000/1688), Loss:loss  
Validation Accuracy of the network: 10.214483014511874 %  
Epoch [4/15], Step (1200/1688), Loss:loss  
Validation Accuracy of the network: 10.19913592574584 %  
Epoch [4/15], Step (1400/1688), Loss:loss  
Validation Accuracy of the network: 10.206045018564357 %  
Epoch [4/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.20267481992797 %



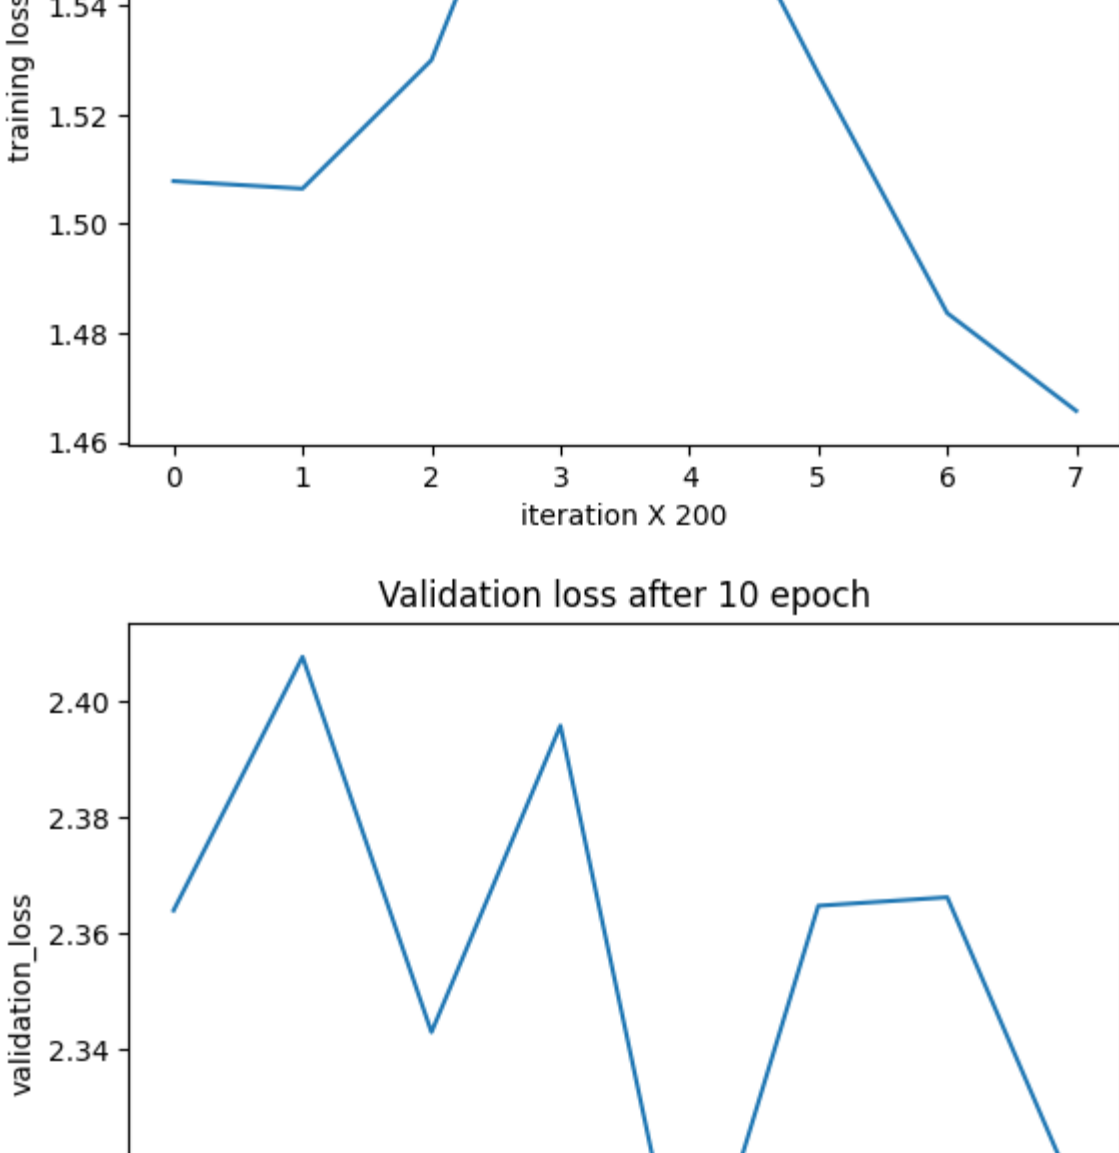

Epoch [5/15], Step (200/1688), Loss:loss  
Validation Accuracy of the network: 10.19086234177251 %  
Epoch [5/15], Step (400/1688), Loss:loss  
Validation Accuracy of the network: 10.185525027964205 %  
Epoch [5/15], Step (600/1688), Loss:loss  
Validation Accuracy of the network: 10.194079842219804 %  
Epoch [5/15], Step (800/1688), Loss:loss  
Validation Accuracy of the network: 10.1843988241525424 %  
Epoch [5/15], Step (1000/1688), Loss:loss  
Validation Accuracy of the network: 10.17721232301341 %  
Epoch [5/15], Step (1200/1688), Loss:loss  
Validation Accuracy of the network: 10.181400905432596 %  
Epoch [5/15], Step (1400/1688), Loss:loss  
Validation Accuracy of the network: 10.16813356849853 %  
Epoch [5/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.168223180076629 %



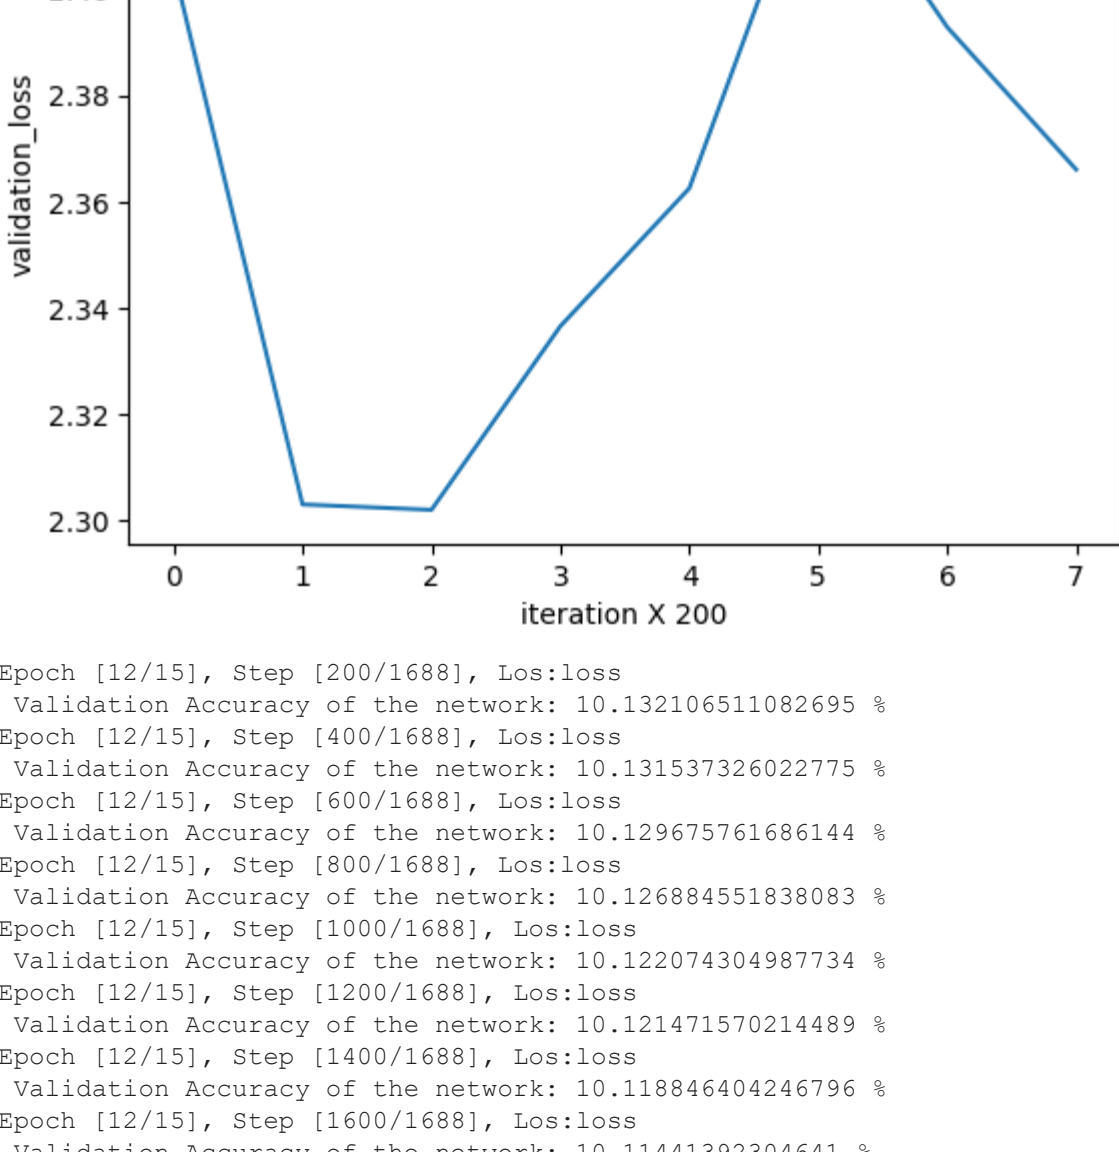
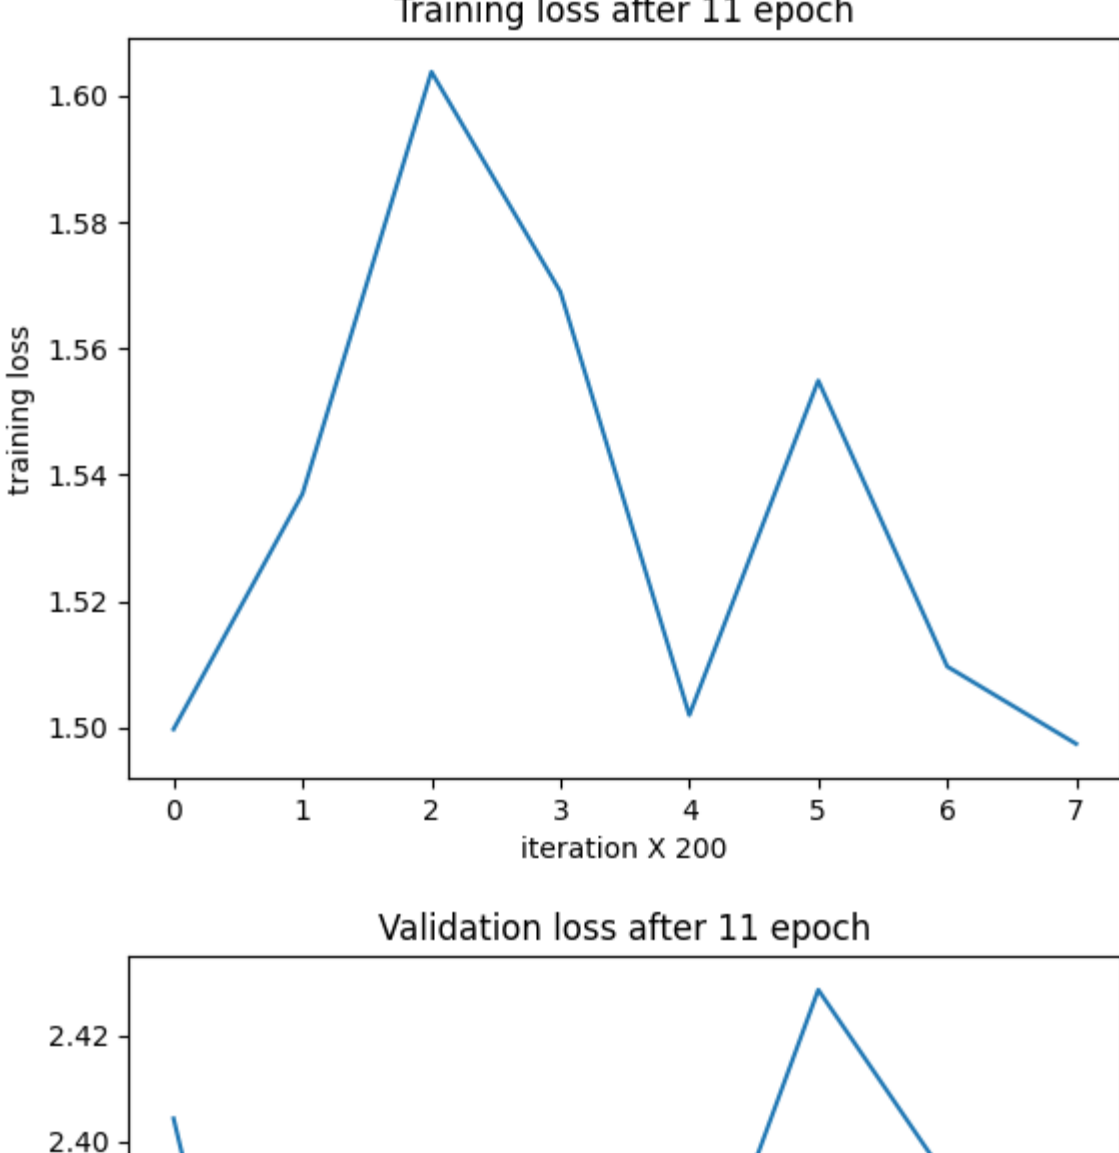
Epoch [6/15], Step (200/1688), Loss:loss  
Validation Accuracy of the network: 10.16285395617196 %  
Epoch [6/15], Step (400/1688), Loss:loss  
Validation Accuracy of the network: 10.17993495475113 %  
Epoch [6/15], Step (600/1688), Loss:loss  
Validation Accuracy of the network: 10.174917035398234 %  
Epoch [6/15], Step (800/1688), Loss:loss  
Validation Accuracy of the network: 10.14229166666666 %  
Epoch [6/15], Step (1000/1688), Loss:loss  
Validation Accuracy of the network: 10.178970593220339 %  
Epoch [6/15], Step (1200/1688), Loss:loss  
Validation Accuracy of the network: 10.177969398340249 %  
Epoch [6/15], Step (1400/1688), Loss:loss  
Validation Accuracy of the network: 10.168953252032521 %  
Epoch [6/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.159673804780876 %



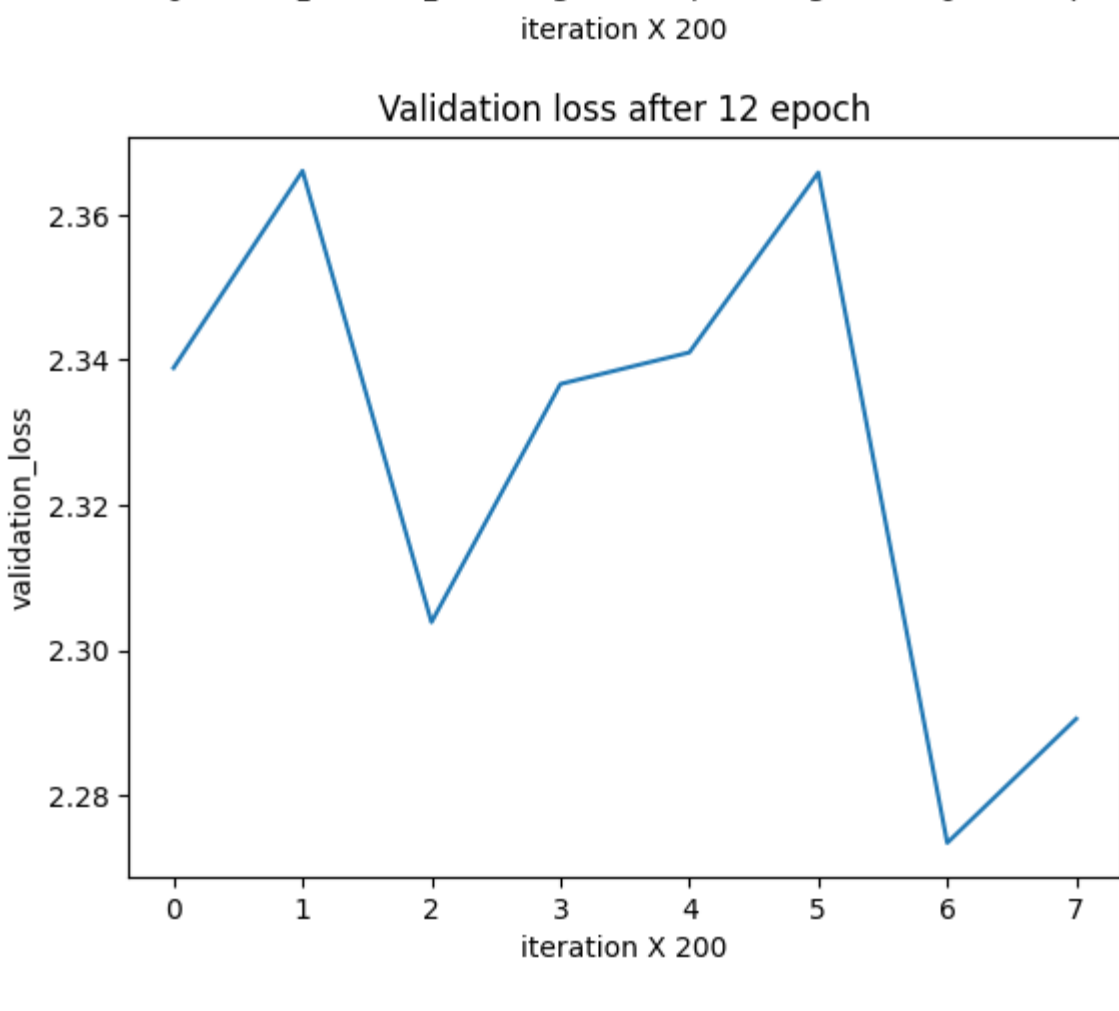
Epoch [7/15], Step (200/1688), Loss:loss  
Validation Accuracy of the network: 10.153829395617196 %  
Epoch [7/15], Step (400/1688), Loss:loss  
Validation Accuracy of the network: 10.162780205167174 %  
Epoch [7/15], Step (600/1688), Loss:loss  
Validation Accuracy of the network: 10.13747426149312 %  
Epoch [7/15], Step (800/1688), Loss:loss  
Validation Accuracy of the network: 10.157997762863534 %  
Epoch [7/15], Step (1000/1688), Loss:loss  
Validation Accuracy of the network: 10.158537701317716 %  
Epoch [7/15], Step (1200/1688), Loss:loss  
Validation Accuracy of the network: 10.1424965708124 %  
Epoch [7/15], Step (1400/1688), Loss:loss  
Validation Accuracy of the network: 10.14990519240113 %  
Epoch [7/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.164942451575295 %  
Epoch [7/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.1610166214188268 %



Epoch [8/15], Step (200/1688), Loss:loss  
Validation Accuracy of the network: 10.13210651082695 %  
Epoch [8/15], Step (400/1688), Loss:loss  
Validation Accuracy of the network: 10.157017436149312 %  
Epoch [8/15], Step (600/1688), Loss:loss  
Validation Accuracy of the network: 10.13747426149312 %  
Epoch [8/15], Step (800/1688), Loss:loss  
Validation Accuracy of the network: 10.142131024096386 %  
Epoch [8/15], Step (1000/1688), Loss:loss  
Validation Accuracy of the network: 10.146685728928784 %  
Epoch [8/15], Step (1200/1688), Loss:loss  
Validation Accuracy of the network: 10.13771847690387 %  
Epoch [8/15], Step (1400/1688), Loss:loss  
Validation Accuracy of the network: 10.145249803097344 %  
Epoch [8/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.164942451575295 %  
Epoch [8/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.13893838214442 %  
Epoch [8/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.134820363744783 %



Epoch [9/15], Step (200/1688), Loss:loss  
Validation Accuracy of the network: 10.129113762405137 %  
Epoch [9/15], Step (400/1688), Loss:loss  
Validation Accuracy of the network: 10.12253668090206 %  
Epoch [9/15], Step (600/1688), Loss:loss  
Validation Accuracy of the network: 10.140842194570136 %  
Epoch [9/15], Step (800/1688), Loss:loss  
Validation Accuracy of the network: 10.128997093023257 %  
Epoch [9/15], Step (1000/1688), Loss:loss  
Validation Accuracy of the network: 10.13964731154362 %  
Epoch [9/15], Step (1200/1688), Loss:loss  
Validation Accuracy of the network: 10.14720077200772 %  
Epoch [9/15], Step (1400/1688), Loss:loss  
Validation Accuracy of the network: 10.13786928726876 %  
Epoch [9/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.152266170155663 %  
Epoch [9/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.15769829184322 %



Epoch [10/15], Step (200/1688), Loss:loss  
Validation Accuracy of the network: 10.15462377962578 %  
Epoch [10/15], Step (400/1688), Loss:loss  
Validation Accuracy of the network: 10.1041655974243 %  
Epoch [10/15], Step (600/1688), Loss:loss  
Validation Accuracy of the network: 10.12253668090206 %  
Epoch [10/15], Step (800/1688), Loss:loss  
Validation Accuracy of the network: 10.146685728928784 %  
Epoch [10/15], Step (1000/1688), Loss:loss  
Validation Accuracy of the network: 10.145249803097344 %  
Epoch [10/15], Step (1200/1688), Loss:loss  
Validation Accuracy of the network: 10.13964731154362 %  
Epoch [10/15], Step (1400/1688), Loss:loss  
Validation Accuracy of the network: 10.13786928726876 %  
Epoch [10/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.152266170155663 %  
Epoch [10/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.134104037363697 %



Epoch [11/15], Step (200/1688), Loss:loss  
Validation Accuracy of the network: 10.1333796837804749 %  
Epoch [11/15], Step (400/1688), Loss:loss  
Validation Accuracy of the network: 10.13747426149312 %  
Epoch [11/15], Step (600/1688), Loss:loss  
Validation Accuracy of the network: 10.140342194570136 %  
Epoch [11/15], Step (800/1688), Loss:loss  
Validation Accuracy of the network: 10.145938199105146 %  
Epoch [11/15], Step (1000/1688), Loss:loss  
Validation Accuracy of the network: 10.121471579214469 %  
Epoch [11/15], Step (1200/1688), Loss:loss  
Validation Accuracy of the network: 10.11884640346796 %  
Epoch [11/15], Step (1400/1688), Loss:loss  
Validation Accuracy of the network: 10.13893838214442 %  
Epoch [11/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.1318930564935 %

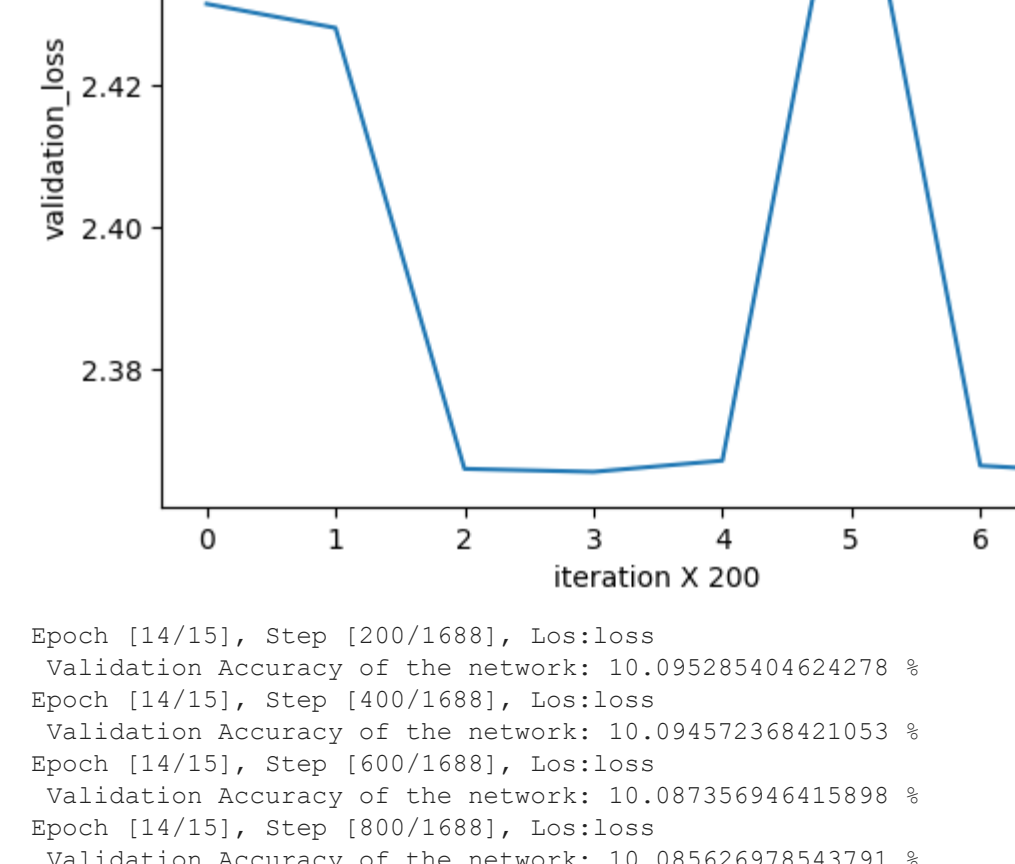



Epoch [12/15], Step (200/1688), Loss:loss  
Validation Accuracy of the network: 10.13210651082695 %  
Epoch [12/15], Step (400/1688), Loss:loss  
Validation Accuracy of the network: 10.13153732602275 %  
Epoch [12/15], Step (600/1688), Loss:loss  
Validation Accuracy of the network: 10.12967561666666 %  
Epoch [12/15], Step (800/1688), Loss:loss  
Validation Accuracy of the network: 10.128997093023257 %  
Epoch [12/15], Step (1000/1688), Loss:loss  
Validation Accuracy of the network: 10.122074304987734 %  
Epoch [12/15], Step (1200/1688), Loss:loss  
Validation Accuracy of the network: 10.121471579214469 %  
Epoch [12/15], Step (1400/1688), Loss:loss  
Validation Accuracy of the network: 10.11884640346796 %  
Epoch [12/15], Step (1600/1688), Loss:loss  
Validation Accuracy of the network: 10.11441392304644 %

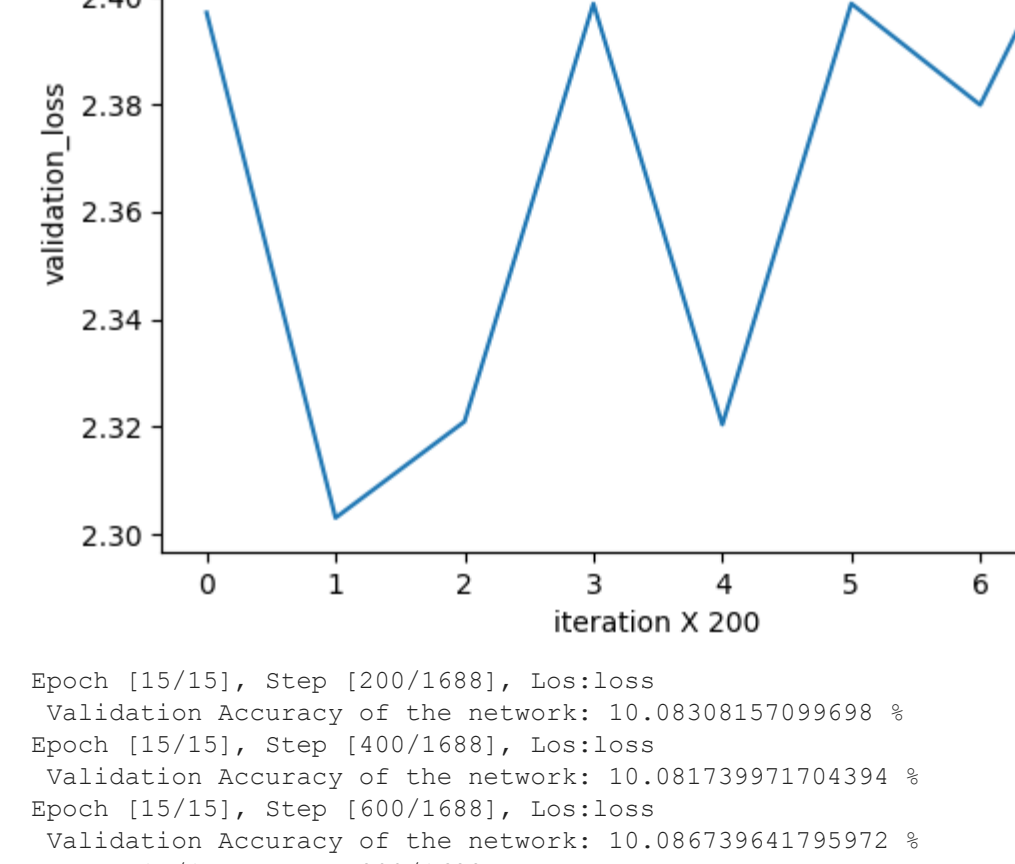
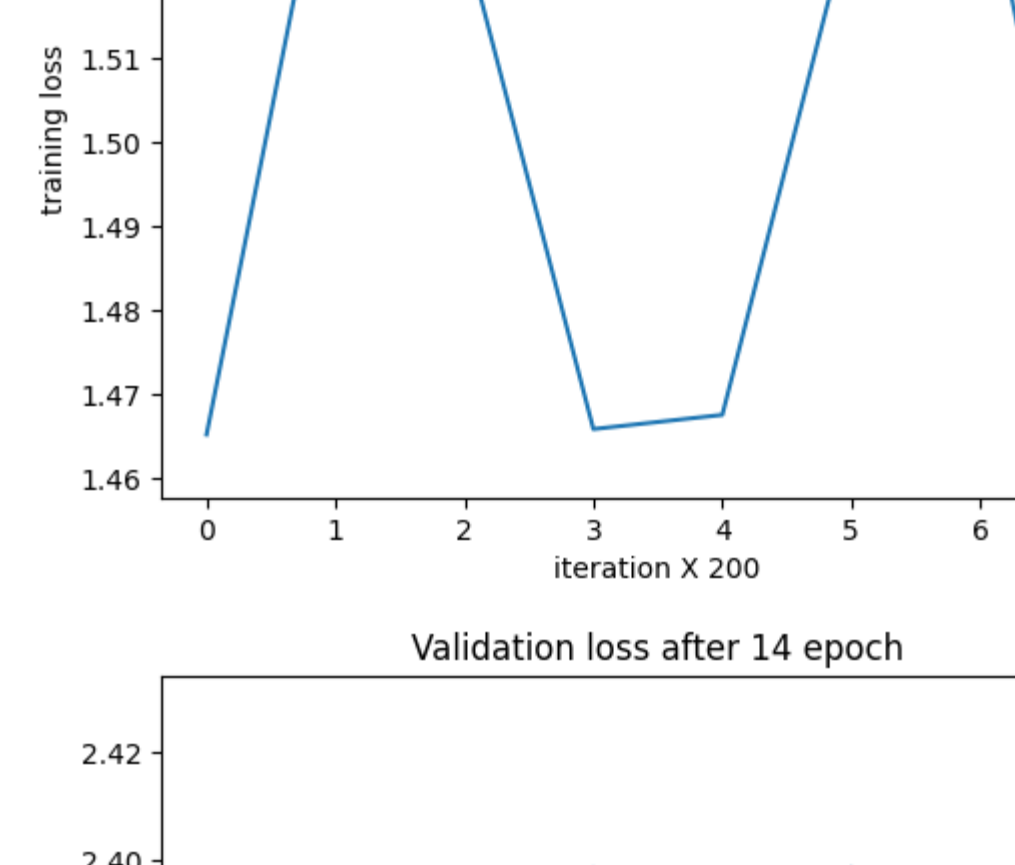







Epoch [13/15], Step [200/1688], Loss:loss  
Validation Accuracy of the network: 10.08892265650645 %  
Epoch [13/15], Step [400/1688], Loss:loss  
Validation Accuracy of the network: 10.109048218435321 %  
Epoch [13/15], Step [600/1688], Loss:loss  
Validation Accuracy of the network: 10.095566024165707 %  
Epoch [13/15], Step [800/1688], Loss:loss  
Validation Accuracy of the network: 10.095973901668912 %  
Epoch [13/15], Step [1000/1688], Loss:loss  
Validation Accuracy of the network: 10.09769994071588 %  
Epoch [13/15], Step [1200/1688], Loss:loss  
Validation Accuracy of the network: 10.08892396028344 %  
Epoch [13/15], Step [1400/1688], Loss:loss  
Validation Accuracy of the network: 10.08409544492178 %  
Epoch [13/15], Step [1600/1688], Loss:loss  
Validation Accuracy of the network: 10.097198709736457 %



Epoch [14/15], Step [200/1688], Loss:loss  
Validation Accuracy of the network: 10.085285404624278 %  
Epoch [14/15], Step [400/1688], Loss:loss  
Validation Accuracy of the network: 10.094572368421053 %  
Epoch [14/15], Step [600/1688], Loss:loss  
Validation Accuracy of the network: 10.08739641795972 %  
Epoch [14/15], Step [800/1688], Loss:loss  
Validation Accuracy of the network: 10.08629694543791 %  
Epoch [14/15], Step [1000/1688], Loss:loss  
Validation Accuracy of the network: 10.088966656572165 %  
Epoch [14/15], Step [1200/1688], Loss:loss  
Validation Accuracy of the network: 10.08728659572165 %  
Epoch [14/15], Step [1400/1688], Loss:loss  
Validation Accuracy of the network: 10.087088526685843 %  
Epoch [14/15], Step [1600/1688], Loss:loss  
Validation Accuracy of the network: 10.082452004975705 %



Epoch [15/15], Step [200/1688], Loss:loss  
Validation Accuracy of the network: 10.08892396028344 %  
Epoch [15/15], Step [400/1688], Loss:loss  
Validation Accuracy of the network: 10.090762115258677 %  
Epoch [15/15], Step [600/1688], Loss:loss  
Validation Accuracy of the network: 10.088629628126014 %  
Epoch [15/15], Step [800/1688], Loss:loss  
Validation Accuracy of the network: 10.08728659572165 %  
Epoch [15/15], Step [1000/1688], Loss:loss  
Validation Accuracy of the network: 10.087088526685843 %  
Epoch [15/15], Step [1200/1688], Loss:loss  
Validation Accuracy of the network: 10.087088526685843 %  
Epoch [15/15], Step [1400/1688], Loss:loss  
Validation Accuracy of the network: 10.087088526685843 %  
Epoch [15/15], Step [1600/1688], Loss:loss  
Validation Accuracy of the network: 10.087088526685843 %



Accuracy of the network: 97.22 %  
Accuracy of 0: 98.28571428571429 %  
Accuracy of 1: 98.9427312753304 %  
Accuracy of 2: 95.25193798449612 %  
Accuracy of 3: 98.41581556415841 %  
Accuracy of 4: 99.18533604887983 %  
Accuracy of 5: 96.6367713044843 %  
Accuracy of 6: 97.2860125560603 %  
Accuracy of 7: 97.37354085603113 %  
Accuracy of 8: 95.1745739876797 %  
Accuracy of 9: 94.4495045495812 %

## Dimensions of the input and output at each layer

output of the filter with kernelize = K X K, stride = S, padding = P for input imagesize = W X W is given by  $\text{floor}(\frac{(W-K+2P)}{S} + 1)$  X  $\text{floor}(\frac{(H-K+2P)}{S} + 1)$

- input MNIST data dimension is 28 \* 28
- so for a batch size of 32 the input dimension for the conv1 layer is 28 X 28 X 32
- conv-1 layer contains the 32 filters with kernelize = 3X3, stride = 1, zero padding = 1 so output size is 28X28X32
- input is 28X28X32, for 2X2 maxpool with stride = 2 has output size of 14X14X32
- input is 14X14X32, for conv-2 layer contains the 32 filters with kernelize = 3X3, stride = 1, zero padding = 1 so output size is 14X14X32
- input is 14X14X32, for 2X2 maxpool with stride = 2 has output size of 7X7X32
- after flattening the outputsize is 7X7X32 = 1568
- input is 1568 and output after FC1 is 500
- input is 500 and output after FC2 is 10

in [23]: `CNN().forward`

Out[23]: `<bound method CNN.forward of CNN(  
 (conv1): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
 (conv2): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
 (fc1): Linear(in_features=1568, out_features=500, bias=True)  
 (fc2): Linear(in_features=500, out_features=10, bias=True)  
 (act1): ReLU()  
>`

## Total number of parameters

in [31]: `summary(CNN(), input_size=(1, 28, 28), device = 'cpu')`

Layer (type)	Output Shape	Param #
Conv2d-1	[1, 32, 28, 28]	320
ReLU-2	[1, 32, 28, 28]	0
Conv2d-3	[1, 32, 14, 14]	9,248
ReLU-4	[1, 32, 14, 14]	0
Linear-5	[1, 500]	784,500
ReLU-6	[1, 500]	0
Linear-7	[1, 10]	5,010

Total params: 799,078  
Trainable params: 799,078  
Non-trainable params: 0  
Params size (MB): 3.05  
Input size (MB): 0.00  
Forward/backward pass size (MB): 0.49  
Params size (MB): 3.05  
Estimated Total Size (MB): 3.54

## NO of parameters in conv and FC layers

- by referring the top of table:
- total no of parameters in convolution layers = 320 + 9,248 = 9,568 (~2% parameters are in convolution layers)
- total no of parameters in fully connected layers = 784,500 + 5,010 = 7,89,510 (almost 98% parameters are contained in fully connected layers) 10\*2078 + 10\*2078

## NO of Neurons

1) The number of neurons in i/p are 28 X 28 = 784 2) The number of neurons after conv1 layer are 32 X 28 X 28 = 25088 3) The number of neurons after maxpool layer are 32 X 14 X 14 = 6272 4) The number of neurons after conv2 layer are 32 X 14 X 14 = 6272 5) The number of neurons after maxpool 2 layer are 32 X 7 X 7 = 1568 6) The number of neurons after FC1 layer are 500 7) The number of neurons after FC 2 layer are 10

- The number of neurons in Convolutional layers = 32\*28\*28 + 6272 + 6272 + 38416
- The number of neurons in FC layers are 1568 + 500 + 10 = 2078

## Batch Normalization

- I have seen that Batch normalization leads to faster convergence and increased accuracy

## Visualization of Conv layers

- please refer this article visualization reference:
- <https://debuggarcia.com/visualizing-filters-and-feature-maps-in-convolutional-neural-networks-using-pytorch>

in [11]: `model_weights = [] # we will save the conv layer weights in this list  
conv_layers = [] # we will save the conv layers in this list  
model_children = list(model.children())`

in [12]: `counter = 0  
# append all the conv layers and their respective weights to the list  
for i in range(len(model_children)):  
 if type(model_children[i]) == nn.Conv2d:  
 counter += 1  
 model_weights.append(model_children[i].weight)  
 conv_layers.append(model_children[i])  
 elif type(model_children[i]) == nn.Sequential:  
 for j in range(len(model_children[i])):  
 if child in model_children[i][j].children():  
 if type(child) == nn.Conv2d:  
 counter += 1  
 model_weights.append(child.weight)  
 conv_layers.append(child)`

print(f"Total convolutional layers: {counter}")  
Total convolutional layers: 2  
in [13]: `for weights, conv in zip(model_weights, conv_layers):  
 # print("WEIGHT: (weight) UNSHAPED: (weight.shape)")  
 print(f"CONV: (conv) ==> SHAPE: (weight.shape)")`


CONV: Conv2d(1, 32, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1)) ==> SHAPE: torch.Size([32, 1, 3, 3])  
CONV: Conv2d(32, 32, kernel\_size=(3, 3), stride=(1, 1), padding=(1, 1)) ==> SHAPE: torch.Size([32, 32, 3, 3])

## 2. Visualizing the CNN

### Conv-1 layer

- by seeing those patterns these are simply searching for the edges and various patterns across the images

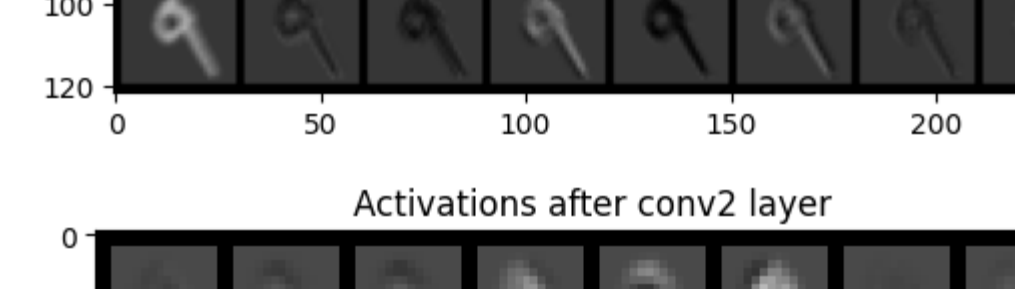
in [14]: `plt.figure(figsize=(20, 17))  
for i, filter in enumerate(model_weights[0]):  
 filter = filter.detach().numpy()  
 plt.subplot(8, 8, i+1) # (8, 8) because in conv0 we have 7x7 filters and total o  
 plt.imshow(filter, cmap='gray')  
 plt.axis('off')`



### Conv-2 layer

- We created 32 channels from the original 1 dimension our second convolution layer adds to this and take 32 channel and outputs 32 different channels. These are second layer channels and are supposed to be more specific.
- Given the size 3X3 for each kernel we can see the image of that size is not big enough for us to visualize what exactly is the model training for but these feature are extracted from the feature which were already highlighted by the 1st layer of convolution.

in [15]: `plt.figure(figsize=(20, 17))  
for i, filter in enumerate(model_weights[1]):  
 filter = filter.detach().numpy()  
 plt.subplot(8, 8, i+1) # (8, 8) because in conv0 we have 7x7 filters and total o  
 plt.imshow(filter, cmap='gray')  
 plt.axis('off')`



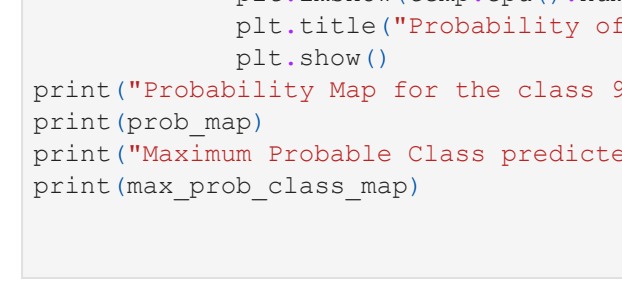
## i have chosen the 20th filter for visualization ,totally 32 filters are available

in [16]: `index = 7  
test_img = test_loader.dataset.data[index].unsqueeze(1).clone()  
test_img = test_img.reshape(1,1,28,28).clone().float()  
ntest_img = test_img[0].permute(1,2,0)  
new_test_img = torch.zeros(28,28)  
new_test_img[0,1:] = ntest_img[1,0]  
print(test_img.size())  
plt.imshow(test_img)  
plt.imshow(new_test_img)  
plt.imshow(test_img.to(device))  
test_img = test_img.to(device)`

in [14]: `torch.no_grad():  
 conv1_out = model.conv1.forward(test_img).reshape(32,1,28,28)  
 conv1_out = conv1_out.cpu()  
 conv1_out = conv1_out - conv1_out.min()  
 conv1_out = conv1_out/conv1_out.max()  
 vis_conv1 = make_grid(conv1_out)  
 plt.imshow(vis_conv1.permute(1,2,0))  
 plt.title("Activations after conv1 layer")  
 plt.show()`

in [15]: `conv2_temp = model.conv1.forward(test_img)  
conv2_temp = model.act1.forward(conv2_temp)  
conv2_temp = model.pool(conv2_temp)  
conv2_out = model.conv2.forward(conv2_temp)  
conv2_out = conv2_out.cpu()  
conv2_out = conv2_out - conv2_out.min()  
conv2_out = conv2_out/conv2_out.max()  
conv2_out = conv2_out.reshape(32,1,14,14)  
vis_conv2 = make_grid(conv2_out)  
plt.imshow(vis_conv2.permute(1,2,0))  
plt.title("Activations after conv2 layer")  
plt.show()`

torch.Size([1, 1, 28, 28])



Activations after conv1 layer



Activations after conv2 layer



- After visualizing the above activationmaps in convlayer-1 and convlayer-2 we can conclude that feature maps become sparse when we going in deeper
- meaning that filter can detect less number of features
- Refer this article <https://towardsdatascience.com/applied-deep-learning-part-4-convolutional-neural-networks-584bc134c1e249722>

## Occulding parts of the image

- by visualizing the below plots and matrix's we can say that learning is meaningful.

in [17]: `dim = len(range(0,14,2))  
max_prob_map = np.zeros((dim, dim), dtype = float)  
max_prob_class_map = np.zeros_like(max_prob_map)  
for y in range(0,14,2):  
 for x in range(0,14,2):  
 temp = test_img.clone()  
 temp[temp>0.9] = 0  
 with torch.no_grad():  
 temp = temp.clone().reshape(1,1,28,28).float()  
 temp = temp.to(device)  
 out = model.forward(temp, softmax = False)  
 probs = F.softmax(out, dim=1).cpu().detach().numpy()  
 pred = np.argmax(probs)  
 prob = probs[:, 9]  
 max_prob_class = pred  
 prob_map[int(y/2), int(x/2)] = prob[0]  
 max_prob_class_map[int(y/2), int(x/2)] = max_prob_class  
 if (x%4 == 0 & y%4 == 0):  
 plt.imshow(temp.cpu().numpy().reshape(28, 28))  
 plt.title("Probability of 9 is 1")*.format(prob[0])  
 plt.show()`

print("Probability Map for the class 9 as positive as patch is moved is shown below:")  
print(prob\_map)  
print("Maximum Probable Class predicted as the patch is moved is shown below:")  
print(max\_prob\_class\_map)



